
Parallel Algorithms and Parallel Computers (ii)

IN4026

Lecture 2

Cees Witteveen

Algorithmics Group

Electrical Engineering, Mathematics and Computer Science

Plan of Lecture 2

- Exercises

- proof of correctness: program cumulative frequencies
- construction program for detecting weakly increasing runs

during lecture

- Prefix computations and balanced trees

- inner product computations
- matrix multiplication

- Pointer jumping

- searching for the root of a tree
- determining the distance to the root

- Divide & Conquer

- finding the minimum 1-index in an array
- finding the maximum of an array

Last week's exercise

Let $A[1..n]$ be an array of n integers.

A is said to contain a *weakly increasing run* of size k if there exists a $j \geq 1$, such that $A[i] \leq A[i+1]$ for $i = j, \dots, j+(k-1)$.

Given such an array A and a number k , decide in $O(\log n)$ parallel time whether A contains a weakly increasing run of at least size k .

Solution

Idea:

1. For $i=1, \dots, n-1$, copy the results of comparing $A[i] \leq A[i+1]$ in an array $B[0..n]$, where $B[0] = B[n] = 0$.
2. Use a balanced tree computation to determine the length of the runs and whether there is a run of at least length k .

To this end we need to define a new binary associative operator $\#$ that can be used to determine lengths of runs.

We use an example to help in defining this operator.

Example

We have an array $A = 1\ 1\ 2\ 4\ 1\ 2\ 1\ 2\ 2\ 3\ 5\ 6\ 4\ 2$
and need to decide whether A contains a run of length $k \geq 5$

1. copy results of comparing $A[i]$ with $A[i+1]$ in an array B

$A = 1\ 1\ 2\ 4\ 1\ 2\ 1\ 2\ 2\ 3\ 5\ 6\ 4\ 2$

$B = 0\ 1\ 1\ 1\ 0\ 1\ 0\ 1\ 1\ 1\ 1\ 1\ 0\ 0\ 0\ 0$ % fill until n is power of 2

Example

We have an array $A = 1\ 1\ 2\ 4\ 1\ 2\ 1\ 2\ 2\ 3\ 5\ 6\ 4\ 2$
and need to decide whether A contains a run of length $k \geq 5$

1. copy results of comparing $A[i]$ with $A[i+1]$ in an array B

$A = 1\ 1\ 2\ 4\ 1\ 2\ 1\ 2\ 2\ 3\ 5\ 6\ 4\ 2$

$B = 0\ 1\ 1\ 1\ 0\ 1\ 0\ 1\ 1\ 1\ 1\ 1\ 0\ 0\ 0\ 0$ % fill until n is power of 2

2. determine run lengths by balanced tree method

$B: 0\ 1\ 1\ 1\ 0\ 1\ 0\ 1\ 1\ 1\ 1\ 1\ 0\ 0\ 0\ 0$
#

$B: (0,1)\ 2\ (0,1)\ (0,1)\ 2\ 2\ 0\ 0$ % keep track of boundaries of runs
#

$B: (0,3)\ (0,1)\ 4\ 0$ % if you see a run of insufficient length,
delete it

$B: (0,1)\ (4,0)$
#

TU Delft $B: true$ % the previous parts form a run of length 6

Solution (i): defining the operator

$0 \# 0 = 0$	$(0,x) \# 0 = \text{true if } x \geq k-1$ $= 0 \text{ else}$	$(x,0) \# 0 = (x,0)$
$0 \# x = (0,x) \text{ if } x > 0$	$(0,x) \# y = (0, x+y) \text{ } y > 0$	$(x,0) \# y = (x,0,y)$
$0 \# (0,x) = (0,x)$	$(0,x) \# (0,y) = \text{true if } x \geq k-1$ $= (0,y) \text{ else}$	$(x,0) \# (y,0) = \text{true if } y \geq k-1$ $= (x,0) \text{ else}$
$0 \# (x,0) = \text{true if } x \geq k-1$ $= 0 \text{ else}$	$(0,x) \# (y,0) = \text{true if } x+y \geq k-1$ $= 0 \text{ else}$	$(x,0) \# (0,y) = (x,0,y)$
$0 \# (x,0,y) = \text{true if } x \geq k-1$ $= (0,y) \text{ else}$	$(0,x) \# (y,0,z) = \text{true if } x+y \geq k-1$ $= (0,z) \text{ else}$	$(x,0) \# (y,0,z) = \text{true if } y \geq k-1$ $= (x,0,z) \text{ else}$
$x \# y = x+y \text{ if } x,y > 0$	$(x,0,y) \# 0 = \text{true if } y \geq k-1$ $= (x,0) \text{ else}$	<ul style="list-style-type: none"> • # true = true true # • = true
$x \# 0 = (x,0) \text{ if } x > 0$	$(x,0,y) \# z = \text{true if } y+z \geq k-1$ $= (x,0,y+z) \text{ else}$	
$x \# (0,y) = (x,0,y) \text{ if } x,y > 0$	$(x,0,y) \# (z,0) = \text{true if } y+z \geq k-1$ $= (x,0) \text{ else}$	
$x \# (y,0) = \text{true if } x+y \geq k-1$ $= (x+y,0) \text{ else}$	$(x,0,y) \# (0,z) = \text{true if } y \geq k-1$ $= (x,0,z) \text{ else}$	
$x \# (y,0,z) = \text{true if } x+y \geq k-1$ $= (x+y,0,z) \text{ else}$	$(x,0,y) \# (u,0,v) = \text{true if } y+u \geq k-1$ $= (x,0,v) \text{ else}$	

Solution (ii): using the balanced tree

Use the following balanced tree scheme:

begin

1. for $1 \leq i \leq n-1$ pardo

if ($A[i] \leq A[i+1]$) then $B[i] := 1$ else $B[i] := 0$;

$B[0] := 0, B[n] := 0$;

2. for $h = 1$ to $\log n + 1$ do

for $0 \leq j \leq ((n + 1)/2^h) - 1$ pardo

$B[j] := B[2j] \# B[2j+1]$;

3. return $B[0]$;

end

$$T(n) = O(\log n)$$
$$W(n) = O(n)$$

Plan of Lecture 2

- Exercises
 - proof of correctness program cumulative frequencies
 - construction program for detecting weakly increasing runs
- Prefix computations and balanced trees
 - inner product
 - matrix multiplication
- Pointer jumping
 - searching for the root of a tree,
 - determining the distance to the root
- Divide & Conquer
 - finding the minimum 1-index in an array
 - finding the maximum of an array

Inner product computations

- Let $\mathbf{u} = [u_i]$ and $\mathbf{v} = [v_j]$ be two $n \times 1$ column vectors.
The inner product $\mathbf{u}^T \mathbf{v}$ is defined as

$$\mathbf{u}^T \mathbf{v} = \sum_{i=1..n} u_i v_i = u_1 \times v_1 + u_2 \times v_2 + \dots + u_n \times v_n$$

Computing the inner product by a balanced tree method:

input: $U[1..n], V[1..n]$ where $n = 2^k$; output: $U^T \cdot V$
begin

1. for $1 \leq i \leq n$ pardo
 $C[i] = U[i] \times V[i];$
2. for $h=1$ to $\log n$ do
 for $1 \leq k \leq n/2^h$ pardo
 $C[k] := C[2k-1] + C[2k];$
3. return $C[1];$

end

$$T(n) = O(\log n), W(n) = O(n)$$

Matrix-vector product

input: $A_{n \times n}, B_{n \times 1}, n = 2^k;$

output: $C_{n \times 1} = A \times B$

begin

1. for $1 \leq i, k \leq n$ pardo

$C'[i, k] := A[i, k] \times B[k];$

2. for $h=1$ to $\log n$ do

for $1 \leq i \leq n, 1 \leq k \leq n/2^h$ pardo

$C'[i, k] := C'[i, 2k-1] + C'[i, 2k];$

3. for $1 \leq i \leq n$ pardo

$C[i] := C'[i, 1];$

end

Remember:

$A.B$ can be written as a vector of inner products $A[i]^T.B$

$$T(n) = O(1), W(n) = O(n^2)$$

$$T(n) = O(\log n), W(n) = O(n^2)$$

$$T(n) = O(1), W(n) = O(n)$$

$$\text{Total: } T(n) = O(\log n), W(n) = O(n^2)$$

Matrix vector product c'td

- On a p-PRAM, the time needed by the balanced tree algorithm is

$$T_p(n) = O(W(n)/p + T(n)) = O(n^2/p + \log n)$$

- This implies that for $p = O(n^2/\log n)$ processors the algorithm is *cost-optimal*.

Matrix product: WT

input: $A_{n \times n}, B_{n \times n}, n = 2^k;$

output: $C_{n \times n} = A \times B$

begin

1. for $1 \leq i, j, k \leq n$ pardo

$C'[i, j, k] := A[i, k] \times B[k, j]$

2. for $h=1$ to $\log n$ do

for $1 \leq i, j \leq n, 1 \leq k \leq n/2^h$ pardo

$C'[i, j, k] := C'[i, j, 2k-1] + C'[i, j, 2k]$

3. for $1 \leq i, j \leq n$ pardo

$C[i, j] := C'[i, j, 1]$

end

Remember:

in $C = A.B$ every element $C[i, j]$ is an inner product $A[i]T.B[j]$

$$T(n) = O(1), \quad W(n) = O(n^3)$$

$$T(n) = O(\log n), \quad W(n) = O(n^3)$$

$$T(n) = O(1), \quad W(n) = O(n^2)$$

$$\text{Total: } T(n) = O(\log n), \quad W(n) = O(n^3)$$

Matrix product: analysis

- Since $T(n) = O(\log n)$, $W(n) = O(n^3)$, for p processors we have

$$T_p(n) = O(n^3/p + \log n)$$

- This implies cost-optimality for $p = O(n^3 / \log n)$ processors

Plan of Lecture 2


- Exercises
 - proof of correctness program cumulative frequencies
 - construction program for detecting weakly increasing runs
- Prefix computations and balanced trees
 - inner product
 - matrix multiplication
- Pointer jumping
 - searching for the root of a tree,
 - determining the distance to the root
- Divide & Conquer
 - finding the minimum 1-index in an array
 - finding the maximum of an array

Pointer jumping

- pointer jumping is a technique suitable for fast access in pointer accessible data structures having the form of a forest of *rooted-directed trees*.

Pointer jumping

- pointer jumping is a technique suitable for fast access in pointer accessible data structures having the form of a forest of *rooted-directed trees*.
- a rooted-directed tree is a directed graph $T = (V, E)$ where
 - V contains a special node r , the root of T ;
 - every node $v \in V - \{r\}$ has out degree 1; (r has out degree 0)
 - for every $v \in V - \{r\}$ there is a *unique* path from v to r
- a *forest* is just a set of trees



Pointer jumping: problem statement

- Given:
 - a forest $F = (V, E)$ where $V = \{1, \dots, n\}$.
 F is represented as an array $P[1 \dots n]$
with $P[i] = j$ iff $(i, j) \in E$, i.e., j is parent of i in a tree of F .
- Question:
 - Find the array $S[1..n]$ such that for every j , $1 \leq j \leq n$,
 $S[j]$ is the root of the tree containing j .
The algorithm should have minimal run time $T(n)$.

Sequential strategy:

Reverse pointers in every tree and perform depth first search. Cost: $O(n)$.

Parallel strategy:

pointer jumping (path doubling): for each node, update the successor of that node by that successor's successor.

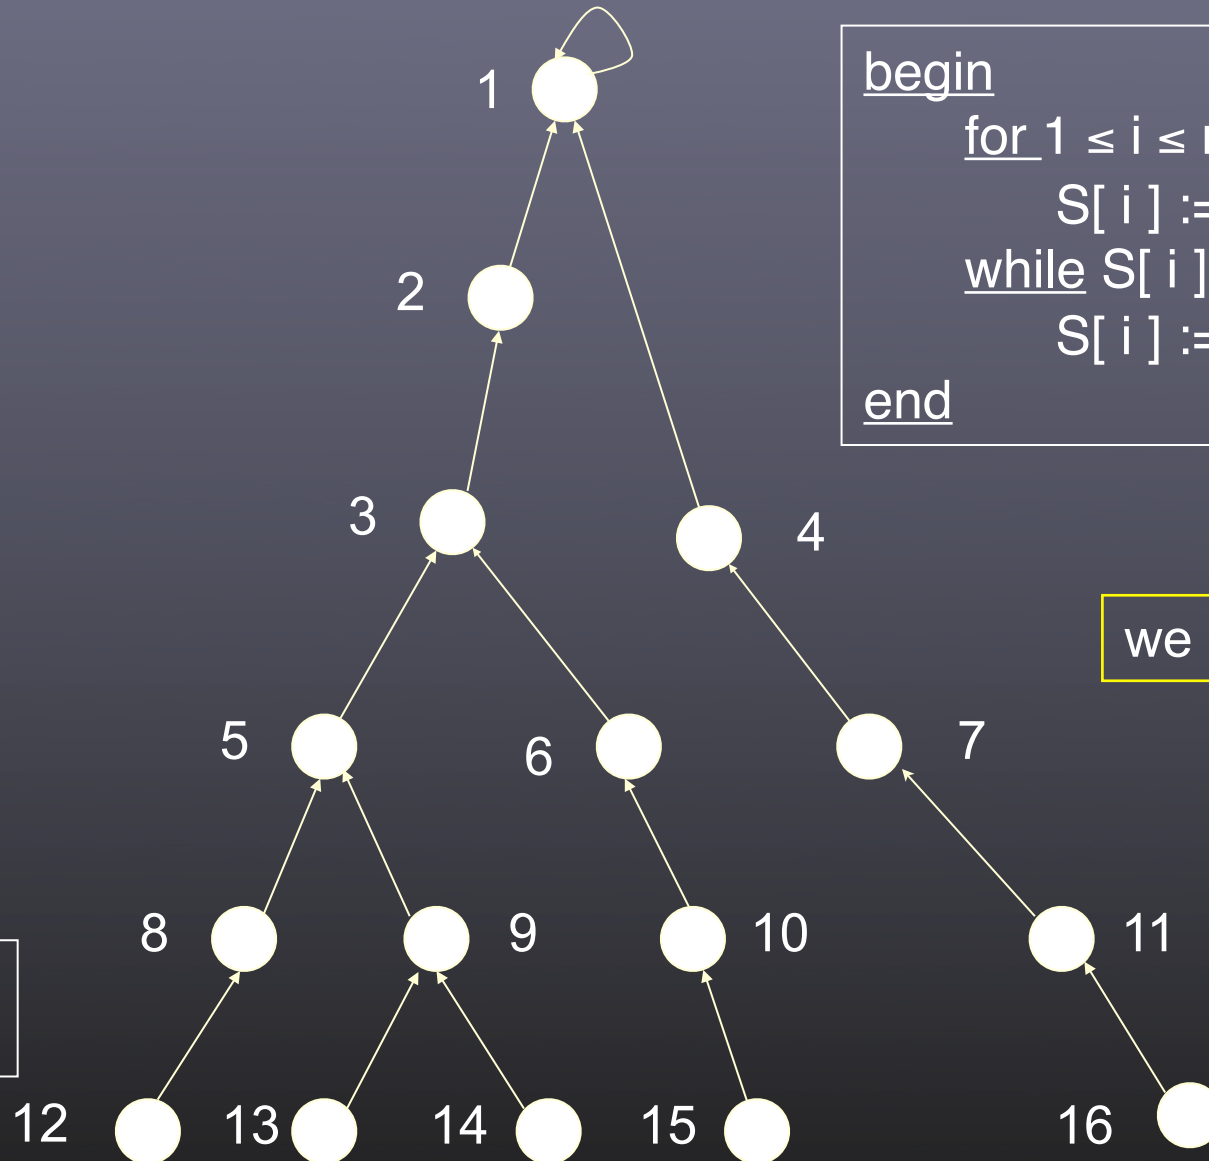
Cost: $T(n) = O(\log h)$, h = height of the tree

Pointer jumping: example

$S = P$

1	1
2	1
3	2
4	1
5	3
6	3
7	4
8	5
9	5
10	6
11	7
12	8
13	9
14	9
15	10
16	11

node - parent
table P



begin

for $1 \leq i \leq n$ pardo

$S[i] := P[i];$

while $S[i] \neq S[S[i]]$ do

$S[i] := S[S[i]];$


end

we assume $P[r] = r$

Pointer jumping: example

S

1	1	1
2	1	1
3	2	1
4	1	1
5	3	2
6	3	2
7	4	1
8	5	3
9	5	3
10	6	3
11	7	4
12	8	5
13	9	5
14	9	5
15	10	6
16	11	7



```

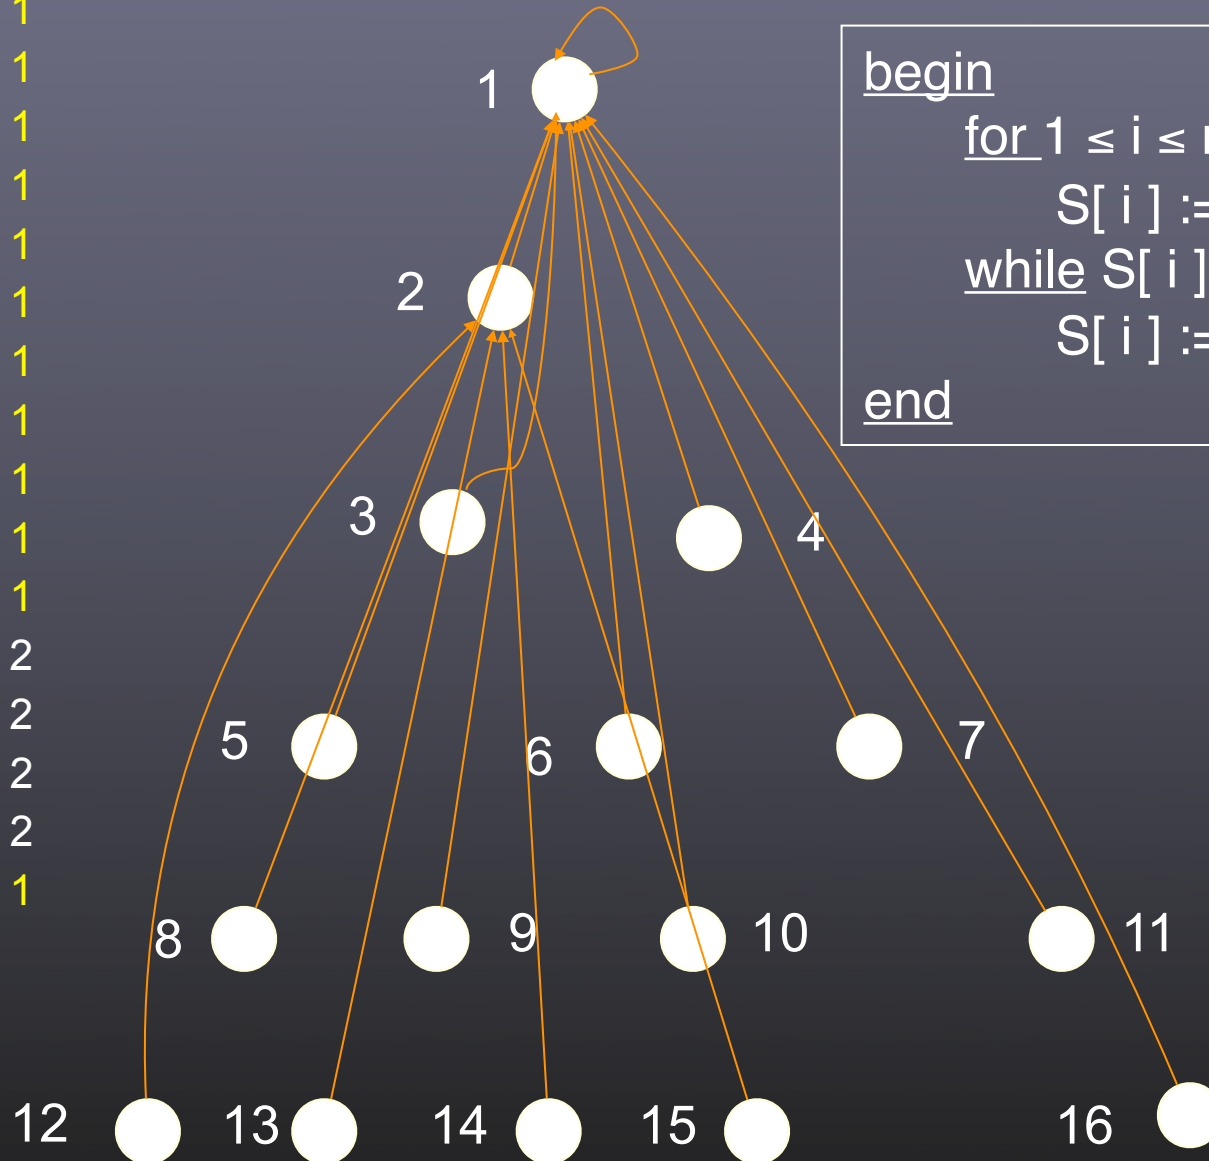
begin
  for  $1 \leq i \leq n$  pardo
     $S[i] := P[i];$ 
    while  $S[i] \neq S[S[i]]$  do
       $S[i] := S[S[i]];$ 
    end
  end

```

Pointer jumping: example

S

1	1	1	1
2	1	1	1
3	2	1	1
4	1	1	1
5	3	2	1
6	3	2	1
7	4	1	1
8	5	3	1
9	5	3	1
10	6	3	1
11	7	4	1
12	8	5	2
13	9	5	2
14	9	5	2
15	10	6	2
16	11	7	1

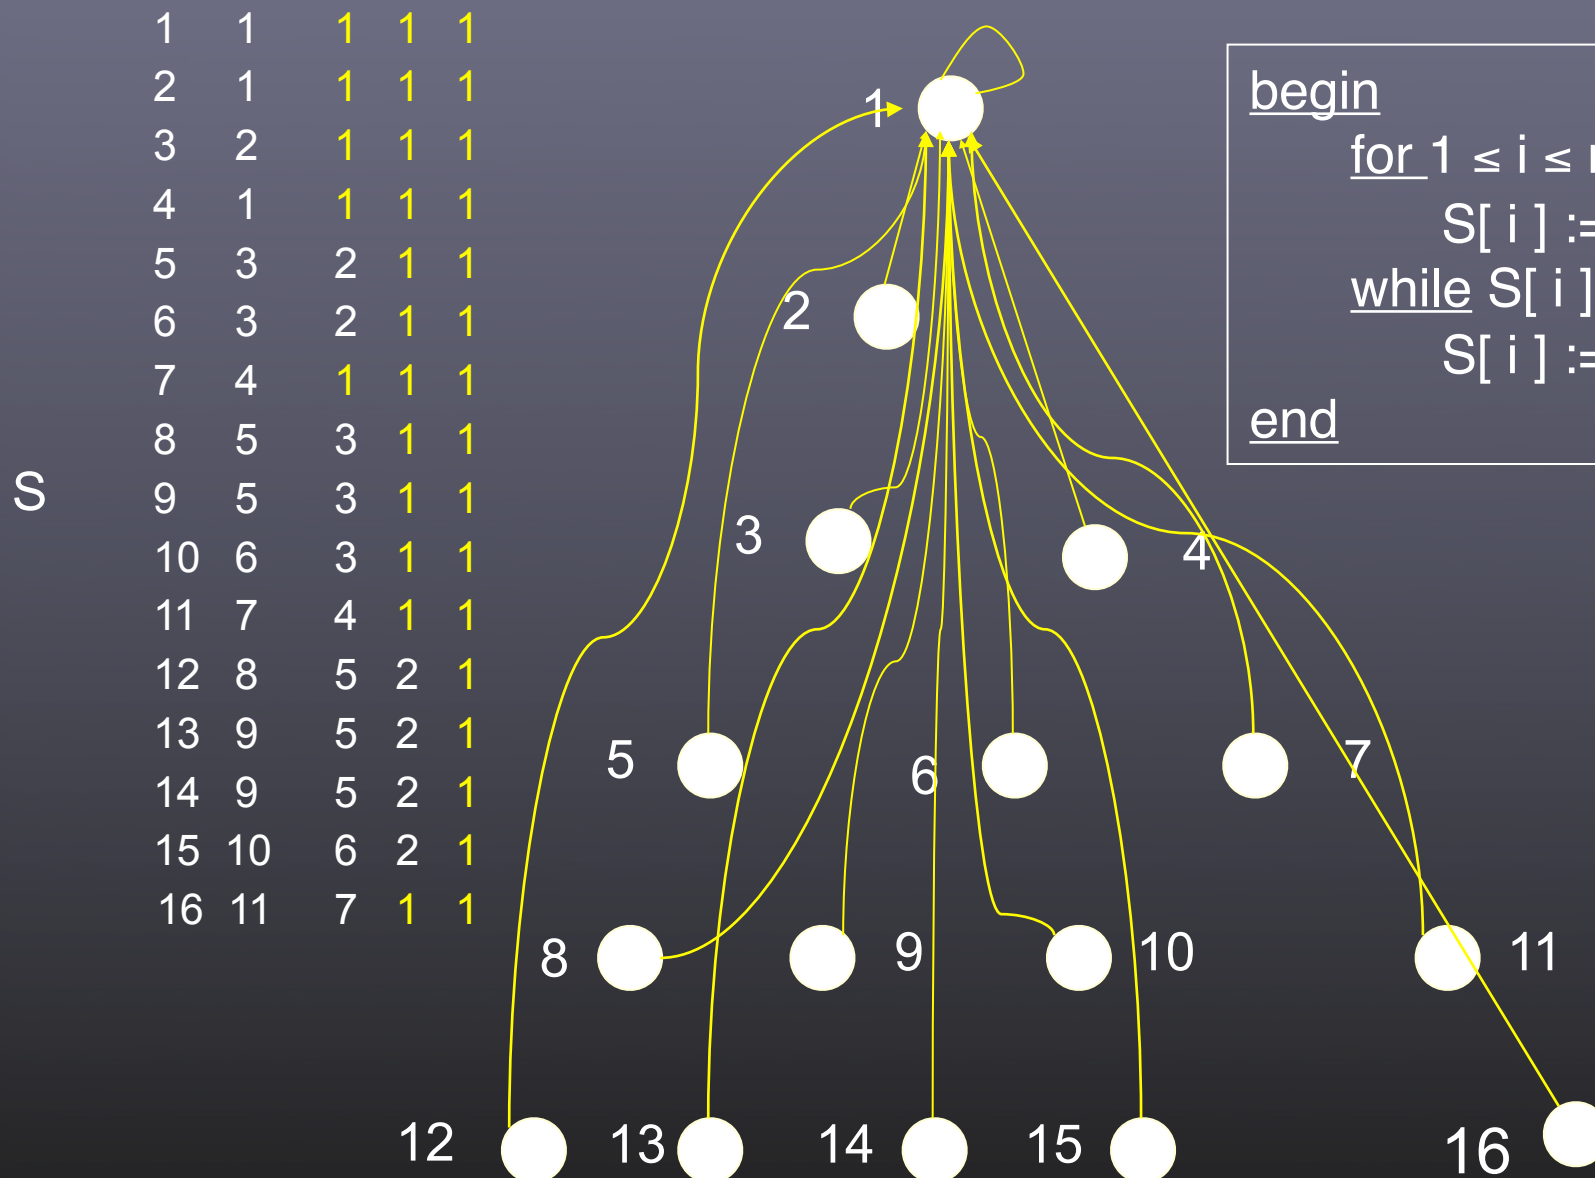


```

begin
  for  $1 \leq i \leq n$  pardo
     $S[i] := P[i];$ 
  while  $S[i] \neq S[S[i]]$  do
     $S[i] := S[S[i]];$ 
  end
end

```

Pointer jumping: example



begin

for $1 \leq i \leq n$ pardo

$S[i] := P[i];$

while $S[i] \neq S[S[i]]$ do

$S[i] := S[S[i]];$

end

Pointer jumping: correctness

input: array P such that $(i, P[i])$ represents edge in E
output: array S with $S[i]$ the root of the tree of node i

begin

for $1 \leq i \leq n$ pardo

$S[i] := P[i];$

$P_0 := S; j := 0$

% dummy statement helping in the proof

while $S[i] \neq S[S[i]]$ do

$S[i] := S[S[i]];$

$j := j+1; P_j := S;$

% dummy statement helping in the proof

end

invariant: $P_k[i] = j$ iff there is a path p from node i to node j
and $((j$ is not root of i and length of p is 2^k) or
 $(j$ is root of i and length of $p \leq 2^k)$).

Pointer jumping: correctness

Since F is a forest, for every node i , there is a unique finite path of length $\leq h$ from i to its root r , where $h = \max \{\text{height of tree in } F\}$.

Consider the following invariant :

$P_k[i] = j$ iff there is a path from node i to node j and $((j$ is not root of i and length of path is 2^k) or $(j$ is root of i and length of path $\leq 2^k)$).

This invariant can be proven correct by induction over $k = 0, 1, \dots, \log h$.

It follows that after $k = \log h$ steps, for every node i , we must have $P_k[i] = P_{k+1}[i]$, that is, $P_k[i]$ contains the root of i .

Therefore, for $k = \log h$, for each i , $S[i] = P_k[i]$ will contain the root of i .

Pointer jumping: analysis

```
begin  
  for  $1 \leq i \leq n$  pardo  
     $S[i] := P[i];$   
    while  $S[i] \neq S[S[i]]$  do  
       $S[i] := S[S[i]];$   
end
```

implementable on
CREW PRAM

Time of algorithm determined by number of iteration steps.
After each iteration step, the distance of node i to node $S(i)$ doubles.
Therefore, we need $\leq \log h$ iterations before $S(i) = r$.
So $T(n) = O(\log h)$.

Every iteration costs $O(n)$ operations.

So $W(n) = O(n \log h)$.

This means that the algorithm is not weakly optimal!

Pointer jumping: other example

Find an algorithm to determine the distances $D[i]$ from node i to the root of the tree:

begin

for $1 \leq i \leq n$ pardo

$S[i] := P[i];$

if $i \neq S[i]$ then $D[i] := 1$ else $D[i] := 0;$

while $S[i] \neq S[S[i]]$ do

$D[i] := D[i] + D[S[i]];$

$S[i] := S[S[i]];$

end

$$T(n) = O(\log h)$$
$$W(n) = O(n \log h)$$

Exercise: use this algorithm to compute the height of a tree.

Special case: linked list

If each tree is just a path represented by a linked list and each node has a weight $W[i]$, we can use pointer jumping to solve the *parallel prefix* problem.

That is, by pointer jumping, we can compute the sum of the weights on the path from node i to its root in $O(\log h)$, $O(n \log h)$ where h is the length of the longest list.

Plan of Lecture 2

- Exercises
 - proof of correctness program cumulative frequencies
 - construction program for detecting weakly increasing runs
- Prefix computations and balanced trees
 - inner product
 - matrix multiplication
- Pointer jumping
 - searching for the root of a tree,
 - determining the distance to the root
- Divide & Conquer
 - finding the minimum 1-index in an array
 - finding the maximum of an array

Divide and Conquer

Basic idea

1. Split problem in nearly equal parts;
2. Solve sub problems concurrently, possibly **recursively** or iteratively,
3. **Combine** solutions of sub problems to solution of the whole problem

Examples of sequential algorithms

binary search;
mergesort;
quicksort

Example mergesort

mergesort (A , n)

if $n=1$ return $A[1]$

else

B := **mergesort**($A[1 \dots n/2]$, $n/2$);

C := **mergesort**($A[n/2+1 \dots n]$, $n/2$);

return merge(B, $n/2$, C, $n/2$)

split and solve separately

combine

merge(X, m, Y, n)

if $m = 1$ and $n=1$ then return $\min(X[1], Y[1]) ++ \max(X[1], Y[1])$

else

if $X[1] \leq Y[1]$ then return $X[1] ++ \text{merge}(X[2 \dots m], m-1, Y, n)$

else

return $Y[1] ++ \text{merge}(X, m, Y[2 \dots n-1], n-1)$

Computing the maximum of an array

Last Lecture's Exercise

Using a CRCW PRAM, compute the maximum of n elements in an array A using an $(O(n^{1+c}), O(1))$ algorithm, where c is an arbitrary positive constant.

Remark:

We already know an $(O(n^{1+c}), O(1))$ algorithm for $c = 1$!


So it is sufficient to look at values $c < 1$.

Solution: the idea

Step 1.

Divide the array of n elements in n^{1-c} blocks of n^c elements.
Compute the n^{1-c} maxima of these blocks in parallel using the $(O(n^2), O(1))$ CRCW-algorithm for computing the maximum.


Solution: the idea



Step 1.

Divide the array of n elements in n^{1-c} blocks of n^c elements. Compute the n^{1-c} maxima of these blocks in parallel using the $(O(n^2), O(1))$ CRCW-algorithm for computing the maximum.


Solution: the idea



Step 2.

Iteratively applying the procedure in **Step 1**, compute the maxima on the resulting arrays with n^{1-c} , $n^{1-c}/n^c = n^{1-2c}$, $n^{1-2c}/n^c = n^{1-3c}$, ..., n^{1-kc} elements containing maxima until $k = \lceil (1-c)/2c \rceil$

Solution: the idea




Step 2.

Iteratively applying the procedure in **Step 1**, compute the maxima on the resulting arrays with n^{1-c} , $n^{1-c}/n^c = n^{1-2c}$, $n^{1-2c}/n^c = n^{1-3c}$, ..., n^{1-kc} elements containing maxima until $k = \lceil (1-c)/2c \rceil$

NB: if $k = \lceil (1-c)/2c \rceil$ then
 $(n^{1-kc})^2 = n^{2-2kc} \leq n^{1+c}$

Solution: the idea




Step 3.

Compute the final maximum for the array containing $\leq n^{1-kc}$ maxima using the $(O(n^2), O(1))$ - algorithm.


Notice that k is a constant depending upon c !

Solution: analysis



$$T = O(1), W = O(n^{1-c} \cdot n^{2c}) = O(n^{1+c})$$

.....



$$T = O(1), W = O(n^{1-(k-1)c} \cdot n^{2c}) = O(n^{1+c})$$



maximum

final maximum

$$T = O(1), W = O(n^{(1-kc)2}) = O(n^{1+c})$$

$$\text{Total: } T(n) = (k+1) \times O(1) = O(1) ; W(n) = (k+1) \times O(n^{1+c}) = O(n^{1+c})$$

Summary of the d&c algorithm

We may assume $c < 1$ (why?).

1. Partition(iteratively) the array in n^{1-c} parts of size n^c .
Apply with the quadratic work-constant time the CRCW PRAM algorithm on every subarray of size n^c and store the result in an array of maxima of size n^{1-c} .
2. Apply step 1 iteratively on the resulting array. The k -th iteration is applied on an array of size $n^{1-(k-1)c}$ and results in an array of size n^{1-kc} . Every iteration costs $T = O(1)$ time and $W = n^{1-kc} \times O(n^{2c}) = O(n^{1+(2-k)c}) \leq O(n^{1+c})$ work.
3. After $k \geq \lceil 1/2c - 0.5 \rceil$ iterations we have $(n^{1-kc})^2 \leq n^{1+c}$.
Hereafter the CRCW-PRAM algorithm can be finally applied on the resulting array with n^{1-kc} maxima to compute the global maximum. The costs are: $T = O(1)$, $W = O(n^{1+c})$.

The total costs are:

$$\begin{aligned} T(n) &= (k+1) O(1) = O(1), \\ W(n) &= (k+1) O(n^{1+c}) = O(n^{1+c}). \end{aligned}$$

New exercise:

Problem: min-1 index

- Given
a boolean array $A[1..n]$
- Question
find an $W(n) = O(n)$, $T(n) = O(1)$ algorithm on a
CRCW-PRAM to compute the first 1 in A ,
that is, compute *the smallest index* k such that $A[k] = 1$.

See you next week